The unrestricted optimum after Conjecture 20: the one-dimensional mechanism, a width theorem, and what can and cannot be proved

17 October 2026

Abstract

We tried to prove that the unrestricted maximum margin equals the paper's prediction, either the interference-free (IF) value $\nu^{1-s}B_p\Gamma_s(p)$ or the price–programme value $\pi_{n,b}(p)B_p\Gamma_s(p)$ of the companion note. The outcome is as follows. (1) A universal proof of the IF value is impossible. At $p = 3$, $b = 2$ we identify the mechanism behind the counterexamples exactly: a *shared-sign, one-dimensional* family whose feature field is a rotating phase times a real exponential sum. Searching this family (exact LP for the weights, outer search over m real slopes) gives new certified networks with $\gamma/\gamma_{\mathrm{IF}} = 1.063, 1.193, 1.184, 1.457, 1.617, 1.595, 2.011, 2.212$ for $n = 11, \dots, 18$. The excess grows with n. (2) We prove a *width theorem*: a slot-symmetric single-feature network whose input vectors share a sign pattern needs at least $\nu - \lfloor \nu/p \rfloor$ neurons. Every network in (1) attains this bound. The saving that drives the counterexamples is therefore at most a $1/p$ fraction of the neurons, and is zero when $p > \nu$, which covers the paper's experimental range. (3) For $p = 5, 7$ the same mechanism, run at its minimal width, loses to IF by a factor that grows with n (efficiency ≤ 0.19 for $p = 5$, $n \le 16$, and decreasing). Releasing it into the full slot-symmetric class did not help at the point we tested ($n = 10$). (4) Decoupling, the remaining step of the price–programme conjecture, is *not* automatic: at $p = 5$ a mean-zero frequency-2 field carried by the same neurons raises a feature's margin by factors 1.2–2.6. In every case we found it still stays $\ge 4\times$ below IF. So the correct universal statement is the price–programme form, and the open content is concentrated in two precise questions for $p \ge 5$.

1 Setting and notation

As in the paper and the companion note: p odd prime, $\rho = e^{2\pi i/p}$, $\theta := \pi/p$, $\nu = n + 1$, $s = \nu/b$. A *single-feature* network has output vectors at frequency 1 only; its margin is $h = \kappa_{\min}(p) \min_x \mathrm{cm}(\Phi(x))$, where $\Phi(x) = \sum_i W_i \prod_t u_{it}(x_t)\rho^{x_t}$ and $\mathrm{cm}(z) = \mathrm{Re}\, z - \cot\theta\, |\mathrm{Im}\, z|$ (companion note, Lemmas 1–2). So $h > 0$ iff every value $\Phi(x)$ lies in the open sector $|\arg z| < \theta$. A network is *slot-symmetric* if $u_{it} = u_i$ for all t; then Φ depends on x only through the count vector $N \in \mathbb{Z}^p_{\ge 0}$, $N_a = \#\{t : x_t = a\}$:

$$\Phi(N) = \sum_{i=1}^{m} W_i \prod_{a \in \mathbb{Z}_p} \big(u_i(a)\rho^a\big)^{N_a}.$$

Efficiency means $\gamma/\gamma_{\mathrm{IF}}$ at $b = 2$ (for single-feature networks at $p \ge 5$, $n \ge 3$ this is the same as the ratio of the single-feature margin to the IF single-feature margin, because $K^*(s, p) = 1$).

2 Shared-sign networks and their fields

Definition 1. A slot-symmetric network has a *shared sign pattern* $\varepsilon \in \{\pm 1\}^p$ if $\varepsilon_a u_i(a) \geq 0$ for all i, a. The *sine pattern* is $\varepsilon_a = +1$ for $0 \leq a < p/2$ and -1 for $p/2 < a < p$, with exponents $e_a := 2a$ resp. $2a - p$, a permutation of $\{0, \ldots, p-1\}$ satisfying $\varepsilon_a \rho^a = e^{i\theta e_a}$. The *one-dimensional (1-D) family* consists of the sine-pattern networks with $u_i(a) = \varepsilon_a e^{\lambda_i e_a}$ (up to scale), $\lambda_i \in \mathbb{R}$.

Lemma 2 (fields)**.** *For a sine-pattern network, $\Phi(N) = e^{i\theta\tau(N)} S(N)$ with $\tau(N) = \sum_a e_a N_a \in \{0, \ldots, (p-1)n\}$ and $S(N) = \sum_i W_i \prod_a |u_i(a)|^{N_a}$, an exponential sum with positive bases. In the 1-D family*

$$\Phi(N) = e^{i\theta\tau} \sum_{i=1}^{m} V_i e^{\lambda_i \tau}, \qquad \tau = \tau(N),$$

so the margin condition is a condition on the $L+1$ values $\tau = 0, \ldots, L$, $L = (p-1)n$, of a rotating phase times a real-exponential sum.

Proof. $u_i(a)\rho^a = |u_i(a)|\varepsilon_a \rho^a = |u_i(a)| e^{i\theta e_a}$; take the product over a with exponents N_a. □

At $p = 3$ the certified counterexample of the earlier notes ($n = 11$, width 8) is, to three digits, a 1-D network: all eight input vectors have the sign pattern $(-, +, +)$ (a translate of the sine pattern), and $\log|u_i(a)| - \log|u_i(0)| = \lambda_i \cdot (0, -1, 1)$ with ratio -1.00 ± 0.03. Re-optimising it inside the 1-D family raises its efficiency from 1.0544 to 1.0631.

3 A width theorem

Theorem 3 (width)**.** *Let p be an odd prime, $n \geq 1$, and let a slot-symmetric single-feature network with a shared sign pattern have positive margin. Then its width is*

$$m \;\geq\; \nu - \left\lfloor \frac{\nu}{p} \right\rfloor.$$

Proof. Positive margin is an open condition, so we may perturb zero entries of the u_i to small values of the prescribed sign. Hence assume all $u_i(a) \neq 0$. Put $d = (p-1)/2$. The map $a \mapsto a + d$ is a cyclic permutation of $\mathbb{Z}_p$ of odd length, so ε cannot alternate along it, and there is a symbol b with $\varepsilon_b = \varepsilon_{b+d}$. Let $a = b + d$ and consider the $n+1$ count vectors $N(k) = (n-k)\mathbf{e}_b + k\mathbf{e}_a$, $k = 0, \ldots, n$. With $r_i = u_i(a)/u_i(b) > 0$,

$$\Phi(N(k)) = C\left(e^{i(\pi-\theta)}\right)^k S(k), \qquad S(k) = \sum_i c_i r_i^k,$$

for constants $C \in \mathbb{C} \setminus \{0\}$, $c_i \in \mathbb{C}$, because $\rho^d = e^{i(\pi-\theta)}$. Positive margin gives $\arg S(k) \in I_k := (\omega_k - \theta, \omega_k + \theta)$ with $\omega_k = -\arg C - (\pi - \theta)k$. Note $\omega_k \equiv -\arg C + \theta k \pmod{\pi}$.

For $\psi \in \mathbb{R}$ let $h_\psi(k) = \mathrm{Im}(e^{-i\psi} S(k)) = \sum_i \mathrm{Im}(e^{-i\psi} c_i)\, r_i^k$. After merging equal bases this is a real exponential sum with at most m terms. By Laguerre's rule of signs it has at most $m - 1$ real zeros unless it vanishes identically, so the sequence $h_\psi(0), \ldots, h_\psi(n)$ has at most $m - 1$ sign changes. Fix a residue k_0 and take $\psi = -\arg C + \theta k_0$. The arc I_k meets the line $\psi + \pi\mathbb{Z}$ iff $\theta(k - k_0) \in (-\theta, \theta) + \pi\mathbb{Z}$, i.e. iff $k \equiv k_0 \pmod{p}$. Call these k *undetermined*. For every other k the sign of $h_\psi(k)$ is fixed by the side of the line on which I_k lies.

- If $k, k+1$ are both determined, $I_{k+1} = I_k - \pi + \theta$; an arc of length 2θ on one side, rotated by $-\pi + \theta$ and still avoiding the line, lies on the other side, so the signs differ.

- If k is undetermined and $0 < k < n$, then $k \pm 1$ are determined ($p \geq 3$), and with I_k centred on the line, I_{k-1} and I_{k+1} are the arcs of length 2θ immediately on either side of it, so $h_\psi(k-1)$ and $h_\psi(k+1)$ again have opposite signs.

So the determined values alternate in sign and there are $|D| - 1$ sign changes, where D is the set of determined k. $h_\psi \not\equiv 0$, since otherwise every $S(k)$ would lie on the line, contradicting $D \neq \emptyset$. Choosing k_0 in a least populated residue class of $\{0, \ldots, n\}$ gives $|D| = n + 1 - \lfloor \nu/p \rfloor$, hence $m - 1 \geq n - \lfloor \nu/p \rfloor$. □

Corollary 4. *(i) For $p > \nu$ a shared-sign single-feature network needs at least ν neurons, as many as the IF block. (ii) For every p, the width saving is at most $\lfloor \nu/p \rfloor$. For networks with equal amplitudes at $b = 2$, the width factor $(\nu/m)^{s-1}$ in the cost comparison with the IF block is therefore at most $\big(\nu/(\nu - \lfloor \nu/p \rfloor)\big)^{s-1} \leq \big(p/(p-1)\big)^{s-1}$: up to $(3/2)^{s-1}$ at $p = 3$, but at most $(5/4)^{s-1}$ at $p = 5$. (iii) Every certified counterexample in Table 1 attains the bound with equality.*

Remark 5. The paper's experiments use $\nu \leq 5$ and $p \geq 11$, so by (i) the mechanism behind the counterexamples has no width advantage anywhere in that range. Width is not the only lever: per unit cost, a flat sign vector has $2^{n/2}$ times the raw output of a cosine input. But numerically this lever alone never sufficed (Section 5).

4 The $p = 3$ optimum: reduction and certified values

Reduction. In the 1-D family, for fixed slopes $\lambda_1, \ldots, \lambda_m$ the cone constraints at $\tau = 0, \ldots, L$ are linear in $(\operatorname{Re} V_i, \operatorname{Im} V_i)$, and the cost amplitude of neuron i is $\sqrt{p/2}\,|V_i|\,Z(\lambda_i)^n$ with $Z(\lambda)^2 = \sum_{k<p} e^{2\lambda k}$. Maximising the margin under a weighted ℓ_1 budget is therefore a second-order cone programme (an LP after polygonal approximation of $|V_i|$). The concave $\ell_{2/\nu}$ cost is handled by reweighting; in practice the optimal amplitudes are nearly equal and uniform weights suffice. The whole problem reduces to choosing m real slopes. We search them with CMA-ES and continuation in n. All values below are then re-verified independently.

Structure of the optimum. Each term sits at the cost-optimal tilt for the stretch of τ it covers. The cost amplitudes agree to within $\sim 10\%$, consecutive terms have nearly antipodal phases ($\approx 172^\circ$ apart at $n = 11$), and the margin is active at most τ. The field is a 10^{-3}–10^{-4} residue of terms up to 10^3 times larger, kept inside the 60° sector at every count class.

n	11	12	13	14	15	16	17	18
width m (bound $\nu - \lfloor \nu/3 \rfloor$)	8 (8)	9 (9)	10 (10)	10 (10)	11 (11)	12 (12)	12 (12)	13 (13)
$\gamma/\gamma_{\mathrm{IF}}$	1.06308	1.19335	1.18360	1.45687	1.61704	1.59546	2.01105	2.21197
previously certified	1.0544	–	1.1091	1.4293	–	–	–	–

Table 1: $p = 3$, $b = 2$. All values are certified by `certify_1d.py`: weights are rounded to ten significant digits, the margin is computed in exact rational arithmetic from the logits over all count vectors (hence over all 3^n inputs), and $\gamma/\gamma_{\mathrm{IF}}$ is enclosed by interval arithmetic.

Hence $\pi_{n,2}(3)\,\nu^{s-1}$ is at least 1.06, 1.19, 1.18, 1.46, 1.62, 1.60, 2.01, 2.21 for $n = 11, \ldots, 18$. The interference-free prediction underestimates the $p = 3$ optimum by a factor that, numerically, grows by about 10% per unit of n.

5 $p \geq 5$: the mechanism loses

Gain/loss accounting. For a 1-D network with nearly equal amplitudes, $\gamma/\gamma_{\mathrm{IF}} = \zeta \cdot 2^{n/2} \cdot (\nu/m)^{s-1}$. Here $2^{n/2}$ is the raw advantage of flat sign vectors over cosines, $(\nu/m)^{s-1}$ the width saving (capped by Theorem 3), and ζ the fraction of the mid-range envelope that survives cancellation. The available gain per unit of n is $\frac{1}{2}\log\frac{2p}{p-1}$: 0.55 at $p=3$, 0.46 at $p=5$, 0.42 at $p=7$. Numerically, $\log\zeta/n \approx -0.48$ at $p=3$ ($n=18$) and ≈ -0.575 at $p=5$ (stable over $n=13$–17). That leaves net rates of about $+0.07$ and -0.12 per unit of n, consistent with the observed growth at $p=3$ and decline at $p=5$.

n	6	7	8	9	10	11	12	13	14	15	16
$p=5$	.192	.179	.166	.147	.132	.127	.126	.120	.108	.097	.093
width (bound)	6 (6)	7 (7)	8 (8)	9 (8)	10 (9)	10 (10)	11 (11)	12 (12)	13 (12)	14 (13)	14 (14)
$p=7$	.126	.107	.089	.076	.066	.063	.058	.049	.044	.035	–

Table 2: Best 1-D efficiency $\gamma/\gamma_{\mathrm{IF}}$ found (slope search with exact LP, continuation in n with insertion and deletion moves and fresh global candidates at every n).

The 1-D restriction is not the bottleneck. At $p=5$, $n=10$, released into the full slot-symmetric single-feature class (arbitrary input vectors, `cone.py`), the 1-D optimum 0.1318 is a local optimum (trial with no noise: 0.1318; with noise: lower). Mixing in positive-input neurons, which keeps the field a function of τ but breaks the shared sign pattern, lowered the efficiency at $p=3$, $n=11$ ($1.063 \to 0.82$).

6 Cross-feature compensation: decoupling is not automatic

At $p \geq 5$ several frequencies coexist, and the margin condition at x is that the vector $(\Phi_l(x))_l$ lies in a polyhedron, not that each $\Phi_l(x)$ lies in its own sector. Concretely, at $p=5$ a frequency-2 value $\Phi_2 \approx 0.25|\Phi_1|e^{-i18^\circ}$ makes Φ_1 admissible at $\arg\Phi_1 = 60^\circ$, the $p=3$ sector. The same neurons can produce such a Φ_2 almost for free: the field is a 10^{-3} residue of much larger terms, so a frequency-2 output component with $\sim 10^{-3}$ relative weight, costing $\sim 10^{-6}$ of the budget, gives a Φ_2 of the right size. It is limited only by geometry: its phase follows a different linear functional of the counts.

We tested this with an exact LP over output vectors spanning frequencies 1 and 2, at the slopes of the best 1-D networks, on all count classes and all four shifts δ. Measured gains over frequency-1 outputs at the same neurons: $\times 1.17$ ($n=8$), $\times 1.36$ ($n=12$), up to $\times 2.6$ at non-optimal widths. With a CMA search over slopes for the joint problem: 0.197 ($n=8$) and 0.173 ($n=12$). The frequency-2 output carries 0.2–3% of the output energy. So the decoupling half of the price–programme conjecture (companion note, Conjecture 20′) cannot be proved feature by feature: one feature can live outside its own sector where another, interference-only, field compensates. In every configuration we tried, the compensated feature still stays $\geq 4\times$ below IF.

7 The number of features of the unrestricted optimum

The paper's target is the *count*: a frequency j is a feature of θ if $A_j(\theta) = |\beta_j(\theta)| > 0$, and the prediction is that every globally margin-optimal network uses $K^*(s,p) = |\operatorname{spt} y^\star|$ features, with amplitudes proportional to a maximiser $y^\star$ of the programme $\Gamma_s(p)$. This is weaker than the value prediction, and it survives the counterexamples.

Proposition 6 (what is unconditional)**.** *Let $\theta^\star$ be a global maximiser of the $L_{2,b}$-margin.*

(a) If $p = 3$, then $\theta^\star$ uses exactly $K^*(s,3) = 1$ feature, for every n and b.

(b) If $\theta^\star$ is interference-free, its amplitudes are proportional to a maximiser of $\Gamma_s(p)$, so it uses exactly $|\operatorname{spt} y^\star|$ features (paper, Theorem 8(1)).

(c) If $b_0 \le b \le \nu$ (interval of the earlier notes, Part II), then $\theta^\star$ is interference-free and uses all $F = K^*(s,p)$ frequencies with equal amplitudes.

(d) For every p, n, b: write $\eta(\theta) := \|A(\theta)\|_{1/s}/(B_p\|\theta\|_{2,b}^\nu) \le 1$ for the clean efficiency and $y = A/\|A\|_{1/s}$ for the clean allocation of $\theta^\star$. Then $\eta(\theta^\star) \ge \pi_{n,b}(p)$ and $\min_{\delta\neq 0}\sum_j y_j\kappa_j(\delta) \ge \frac{\pi_{n,b}(p)}{\eta(\theta^\star)}\,\Gamma_s(p)$. In particular, if $\eta(\theta^\star) = \pi_{n,b}(p)$, then y is a programme maximiser and $\theta^\star$ uses exactly $|\operatorname{spt} y|$ features.

Proof. (a) $F = \{1\}$, and $A_1 > 0$ by the paper's Lemma 2 because the margin is positive. (b), (c) are the cited results. (d) By the price–programme theorem of the companion note, $h(\theta^\star) = \gamma_b^*\|\theta^\star\|^\nu \ge \pi B_p\Gamma_s\|\theta^\star\|^\nu$. By the paper's Lemma 2, $h \le \min_\delta \sum_j A_j\kappa_j(\delta) = \|A\|_{1/s}\min_\delta\sum_j y_j\kappa_j(\delta) \le \|A\|_{1/s}\Gamma_s$. Divide by $B_p\|\theta^\star\|^\nu$. □

Proposition 7 (decoupling suffices; the price is irrelevant)**.** *Suppose a global maximiser $\theta^\star$ satisfies the decoupling inequality*

$$h(\theta^\star) \;\le\; \pi_{n,b}(p)\,B_p\,\min_{\delta\neq 0}\sum_j P_j^s\kappa_j(\delta) \qquad \text{for some } P_j \ge 0,\;\; \sum_j P_j \le \|\theta^\star\|_{2,b}^b.$$

Then $(P_j^s/\|\theta^\star\|_{2,b}^\nu)_j$ is a maximiser of $\Gamma_s(p)$. If moreover the features of $\theta^\star$ are exactly the frequencies it pays for ($A_j > 0 \iff P_j > 0$), then $\theta^\star$ uses exactly K^* features.

Proof. With $z_j = P_j^s/\|\theta^\star\|^\nu$ we have $\sum_j z_j^{1/s} \le 1$, and

$$\pi B_p\Gamma_s\|\theta^\star\|^\nu \le h(\theta^\star) \le \pi B_p\|\theta^\star\|^\nu\min_\delta\sum_j z_j\kappa_j(\delta).$$

Hence z attains Γ_s. □

The price enters only as a common factor, because relabelling $a \mapsto ja$ makes it the same for every frequency and scaling makes it homogeneous of degree s in the budget. So the count needs *only* decoupling. Neither $\pi = \nu^{1-s}$ nor interference-freeness is needed.

Theorem 8 (the count is $\Theta(\log p)$ for every near-optimal network)**.** *Let $2 \le b < \nu$, $s = \nu/b$, $r = 1/s$, and for a non-zero amplitude vector A let*

$$K_r(A) := \frac{\big(\sum_j A_j^r\big)^{1/(1-r)}}{\big(\sum_j A_j\big)^{r/(1-r)}} \;\le\; |\operatorname{spt} A|,$$

the participation number of A (it equals K for K equal non-zero entries). Let θ be any product network, of any width and with interference allowed, with $\gamma(\theta) \ge \lambda\,\nu^{1-s}B_p\Gamma_s(p)$ for some $\lambda \in (0,1]$. Every global maximiser has $\lambda \ge \pi_{n,b}(p)\nu^{s-1} \ge 1$. Then

$$K_r(A(\theta)) \;\le\; \nu\Big(\frac{p}{(p-1)\lambda\,\Gamma_s(p)}\Big)^{\frac{1}{s-1}} \;\le\; \nu\Big(\frac{2p}{(p-1)\lambda}\Big)^{\frac{1}{s-1}}\lceil 8\ln p\rceil.$$

Moreover, for every $M \geq 1$ the M largest amplitudes carry at least a fraction $1-(K_r/M)^{1-r}$ of $\sum_j A_j$. Together with the paper's Corollary 12, which shows that the support of such a network has at least $c_{\nu,b,\lambda} \ln p$ frequencies, the feature count of every margin-optimal network is $\Theta_{\nu,b}(\log p)$. This is the order of $K^*(s,p)$ (paper, Theorem 11), for every odd prime, unconditionally.

Proof. Averaging the paper's Lemma 2 over $\delta \neq 0$ and using $\sum_{\delta\neq 0} \kappa_j(\delta) = p$ gives $h \leq \frac{p}{p-1} \sum_j A_j$. The single-neuron budget gives $\sum_j A_j^r \leq \sum_{i,j} |\beta_{ij}|^r \leq C_b P$ with $P = \|\theta\|_{2,b}^b$. Hence, with $h = \gamma P^s$ and $sr = 1$,

$$K_r(A) \leq \frac{(C_b P)^{1/(1-r)}}{\left(\frac{p-1}{p}\gamma P^s\right)^{r/(1-r)}} = \left(\frac{p\, C_b^{1/r}}{(p-1)\gamma}\right)^{\frac{r}{1-r}} = \left(\frac{p\, B_p}{(p-1)\gamma}\right)^{\frac{1}{s-1}},$$

because $C_b^{1/r} = C_b^{\nu/b} = B_p$ and $r/(1-r) = 1/(s-1)$. Insert $\gamma \geq \lambda \nu^{1-s} B_p \Gamma_s$, then $\Gamma_s \geq \frac{1}{2}\lceil 8 \ln p \rceil^{1-s}$ (paper, Theorem 10). The bound $\lambda \geq \pi \nu^{s-1}$ for maximisers is the price–programme theorem. For the tail statement, normalise $\sum_j A_j = 1$ and sort; the amplitudes beyond the M-th are at most $1/M$ each, so their mass t satisfies $tM^{1-r} \leq \sum_{j>M} A_j^r \leq K_r^{1-r}$. □

Corollary 9 ($n = 2$, every prime $p \geq 5$)**.** *For $n = 2$ and $2 \leq b \leq 3$, no single-frequency network (inputs in $\operatorname{span}\{\mathbf{1}, \cos, \sin\}$ of one frequency, any width) beats the interference-free single-feature value at any prime $p \geq 5$.*

Proof. For $n = 2$ the robust constant of the companion note's Theorem 14 is $k^*(2,s) = \sup(\operatorname{Re} A - 3^{1-s}(\sum |\beta_i|^{1/s})^s)/(R/\sqrt{2})$ with $R^2 = |B|^2 + |C|^2 + |D|^2$, where B, C, D are the three sign-flipped sums. The four-moment inequality of the earlier notes (Part I, Theorem A: $|A| - R/(2\sqrt{3}) \leq 3^{1-s}(\sum |\beta_i|^{1/s})^s$ for $1 \leq s \leq 3/2$, valid for arbitrary unit phases) gives $k^*(2,s) \leq \sqrt{2}/(2\sqrt{3}) = 1/\sqrt{6} < 0.41 < k_5 = 0.946 \leq k_p$. The companion note's Theorem 14 then applies. □

Lemma 10 (slack identity)**.** *For every logit table (any network, any p, n, any inputs), with margin h and slack $\sigma(x,\delta) := f(x)[y(x)] - f(x)[y(x)-\delta] - h \geq 0$,*

$$\Phi_j(x) = \tfrac{2}{p}\Big(h - \sum_{\delta\neq 0} \sigma(x,\delta)\rho^{j\delta}\Big) \qquad (j \in F).$$

Consequently every sign-flipped coefficient satisfies $|\beta_{j,T}| \leq \frac{2}{p} S$ with $S := \mathbb{E}_x \sum_\delta \sigma = p \sum_{j\in F} \operatorname{Re}\beta_j - (p-1)h$. The same slack S bounds the interference of every *frequency.*

Proof. $\hat{f}_x(j) = \sum_\delta f(x)[y-\delta]\rho^{-j(y-\delta)} = -\rho^{-jy}\sum_{\delta\neq 0}(h + \sigma(x,\delta))\rho^{j\delta}$ because $\sum_\delta \rho^{j\delta} = 0$, and $-\sum_{\delta\neq 0}\rho^{j\delta} = 1$. Then $\beta_{j,T} = \overline{\mathbb{E}_x \Phi_j \chi}$ for a unimodular character χ, and $|\mathbb{E}[\sigma\chi]| \leq \mathbb{E}\sigma$. For S, use $\mathbb{E}_x \Phi_j = \overline{\beta_j}$ and that $\sum_{j\in\mathbb{Z}_p} \hat{\bar{\sigma}}(j) = p\,\bar{\sigma}(0) = 0$. □

Proposition 11 ($n = 2$, any inputs)**.** *For $n = 2$, $2 \leq b \leq 3$, every prime p and every network, $h \leq \frac{p}{p-1}\big(\sum_j \operatorname{Re}\beta_j - \max_j R_j/(2\sqrt{3})\big)$, where $R_j^2 = \sum_{T\neq\emptyset} |\beta_{j,T}|^2$. For single-feature networks (any inputs) this gives $\gamma \leq \frac{p}{p-1}\nu^{1-s} B_p$, i.e. $\pi_{2,b}(p) \leq \nu^{1-s} \cdot \frac{p}{(p-1)\kappa_{\min}(p)}$. This is Part I's sharp result at $p = 3$, but it is weaker than the IF value at $p \geq 5$.*

Proof. Lemma 10 with three non-empty T gives $R_j \leq 2\sqrt{3}(\sum_{j'} \operatorname{Re}\beta_{j'} - \frac{p-1}{p}h)$. For one feature, apply Part I's Theorem A ($|A| - R/(2\sqrt{3}) \leq 3^{1-s}(\sum |\beta_i|^{1/s})^s$; Lemma 6 of the paper supplies the unit phases for arbitrary inputs) and the single-neuron budget. □

Two losses separate Proposition 11 from the exact answer at $p \geq 5$, and both are the multi-frequency obstacle. (i) A single feature's gaps at its non-binding offsets carry structural slack, which uses up S without any interference. (ii) With several frequencies, one slack budget bounds each R_j but not $\sum_j R_j$, so only one frequency's interference is charged. Part II of the earlier notes overcomes (ii) only near $b = \nu$, by mixing the chain matrices of all frequencies.

The general $n = 2$ case (inputs with other frequencies) is also where the L^2 route stops. Input energy ε outside ± 1 costs only $O(\varepsilon)$ of clean amplitude, but it can create non-circular (partner-mode) interference of size $O(\sqrt{\varepsilon})$. So every L^2/moment bound fails in a thin regime next to the interference-free point, and a pointwise argument, like Part I's at $p = 3$, is required.

Numerical tests of the count with general networks. (i) *Width.* For slot-symmetric networks with arbitrary input vectors (any sign patterns) and outputs on all frequencies, we searched for the smallest width with positive margin (elimination from feasible starts, exact LP for the outputs). It equals $\nu - \lfloor \nu/p \rfloor$ in every case tested: $p = 3$, $n = 3, 4, 5, 6, 8$ (widths $3, 4, 4, 5, 6$), and $p = 5$, $n = 4, 5, 6, 7$ (widths $4, 5, 6, 7$), with one output frequency or all of them. So mixing frequencies does not save neurons, and Theorem 3 appears to hold for all slot-symmetric networks. (ii) *Allocation.* At $(p, n, b) = (7, 3, 2)$ the programme has $K^* = 2$ with unequal shares $(0.328, 0, 0.182)$. Global search (inputs by CMA-ES from interference-free blocks on every frequency subset plus perturbations, outputs by exact conic programme with ℓ_r reweighting) gave: one frequency 0.801; any two frequencies 1.00000 with amplitude shares equal to the programme's up to relabelling, unchanged under perturbation; all three frequencies 0.827. At $(5, 2, 2)$ ($K^* = 2$), random-start search reached 0.755 (width 4, one feature), 0.821, 0.941, 0.970 (widths 5–7, two features with near-equal shares), and never exceeded the K^*-feature block.

Where a proof stops. Combining (a)–(c): the count prediction holds for every global optimiser, *except possibly* when $p \geq 5$, $b < b_0$ and the optimiser carries interference. None of the known counterexamples is in that set, since they all have $p = 3$. Closing the remaining case requires a price-sharp upper bound, i.e. Conjecture $20'$ or an equivalent. The only unconditional upper bound, Theorem 4 of the paper, is too weak: it loses the factor $\pi\nu^{s-1}$, which swamps the differences between allocations (removing one feature changes Γ_s by a factor $(1 - w_j)^s$ only, paper Lemma 19). The compensation of Section 6 is the precise danger. It uses frequencies whose clean amplitude is tiny but non-zero and that the programme does not pay for. An optimiser built that way would *over*count, so the strict count statement at $p \geq 5$ is exactly as hard as excluding compensation at the optimum. In every experiment at $p = 5, 7$, compensated networks stay $\geq 4\times$ below the K^*-feature IF block.

8 What is proved, what is false, what is open

Statement	Status
$\gamma_b^* = \nu^{1-s} B_p \Gamma_s(p)$ for all p, n, b (Conjecture 20)	**false**: $p = 3$, $b = 2$, $n \geq 11$; excess $2.21\times$ certified at $n = 18$, growing
$\nu^{1-s} B_p \Gamma_s \leq \pi_{n,b}(p) B_p \Gamma_s \leq \gamma_b^* \leq B_p \Gamma_s$	proved (companion note)
$\gamma_b^* = \pi_{n,b}(3)\, B_3 \Gamma_s(3)$ at $p = 3$	proved (one frequency)
width $\geq \nu - \lfloor \nu/p \rfloor$ for shared-sign networks	proved (Theorem 3)
single-frequency networks cannot beat IF when $k_p \geq k^*(n, s)$	proved (companion note); k^* numerical
$\gamma_b^* = \pi_{n,b}(p) B_p \Gamma_s(p)$ (Conjecture 20′, decoupling)	open; genuine obstacle exhibited (Section 6)
$\pi_{n,b}(p) = \nu^{1-s}$ for $p \geq 5$	open; every mechanism tested loses, by a factor growing in n
feature count of every global optimiser $= K^*(s, p)$, $p = 3$	proved (Prop. 6(a))
feature count $= K^*(s, p)$, all p, $b_0 \leq b \leq \nu$	proved (Prop. 6(c))
feature count $= K^*(s, p)$, all p, given decoupling	proved (Prop. 7); price irrelevant
feature count $= \Theta(\log p) = \Theta(K^*)$, all p, unconditional	proved (Thm. 8 with paper's Cor. 12)
single-frequency networks, $n = 2$, all $p \geq 5$, $b \in [2, 3]$: no gain over IF	proved (Cor. 9)
feature count $= K^*(s, p)$, $p \geq 5$, $b < b_0$	open; equivalent in difficulty to excluding compensation

Answer to "prove that the unrestricted optimum matches the prediction". In the interference-free form this cannot be done: the statement is false at $p = 3$, and increasingly so as n grows. In the price–programme form it holds at $p = 3$ by definition, and there the object to compute is $\pi_{n,b}(3)$. We have shown that, in the class that contains every counterexample, this reduces to choosing $\nu - \lfloor \nu/3 \rfloor$ real slopes, with the weights given by an LP. For $p \geq 5$ the full statement would follow from two claims:

(Q1) $\pi_{n,b}(p) = \nu^{1-s}$, i.e. a single feature cannot be bought below the IF price. Theorem 3 limits the width saving of the $p = 3$ mechanism to $\lfloor \nu/p \rfloor$ neurons, and to none for $p > \nu$. The circularity theorem of the companion note handles single-frequency inputs when $k_p \geq k^*$. A proof needs a cost-weighted (robust) form of the width theorem for arbitrary sign patterns.

(Q2) Decoupling. It has to control compensation by interference-only fields, which demonstrably raises a feature's margin.

Both hold in all our experiments at $p = 5, 7$, by margins of $4\times$ or more that grow with n.

Files. `certify_1d.py` (rigorous certificate for the 1-D family, any p), `oned_certified/` (parameters of the networks in Table 1), `mechanism_scripts/` (the field model, the LP, the slope search and continuation, the joint two-frequency LP, the release test).